

Copilot-Authored Commits: Agent Change Attribution

Problem

When the Copilot agent edits files via tools (`create_file`, `replace_string_in_file`, `apply_patch`, etc.), those changes are written directly to disk with no distinction from developer-made edits. This creates several issues:

- **No attribution** — `git log` shows the developer as the author of AI-generated code
- **No auditability** — impossible to tell which changes were human vs. machine after the fact
- **No traceability** — no record of which model or tools produced the code
- **Compliance risk** — organizations increasingly require disclosure of AI-generated code in version control

GitHub Copilot's own coding agent solves this by committing as `copilot[bot]`, but that's a server-side bot account unavailable to third-party integrations.

Constraints

- **No official Copilot email** — GitHub's `copilot[bot]` is server-side only; no public email exists for use in local `git commit --author`
- Creating a dedicated GitHub machine user costs a seat and requires org admin setup
- GitHub Apps can author commits via API but not through local `git commit`
- Git's `--author` flag requires Name <email> format
- IntelliJ's `CheckinHandler.beforeCheckin()` can only return COMMIT or CANCEL — it cannot modify the commit author or message

Solution

Two git-native mechanisms that require no GitHub account setup and are parseable by CI/tooling:

1. `--author` flag — separates who wrote vs. who approved

Git natively separates **author** (who wrote the code) from **committer** (who approved it). We set the author to Copilot, leaving the committer as the developer.

For the email, we use `copilot@copilot.example` — a reserved domain under RFC 6761 guaranteed to never be a real address.

```
git commit --author="GitHub Copilot (gpt-4.1) <copilot@copilot.example>" -m "Add auth endpoint"
```

```
Author: GitHub Copilot (gpt-4.1) <copilot@copilot.example>
Commit: John Doan <john.doan@citi.com>
```

This gives immediate separation — `git log --author="GitHub Copilot"` returns only AI-generated commits, while the committer field preserves accountability.

2. Generated-by trailer — structured metadata for tooling


Git trailers are key-value metadata at the end of a commit message. GitHub renders them in the commit detail view, and CI pipelines can parse them.


```
Add auth endpoint
```

```
Generated-by: github-copilot
```

End-to-end flow

There are three commit paths. All agent paths produce the same attributed commit format.





Auto-inject attribution

The plugin enforces correct attribution regardless of which commit path the developer uses. There is no "commit anyway under your name" escape hatch.

How it works

`CopilotCheckinHandlerFactory` hooks into IntelliJ's standard commit flow via the `CheckinHandler` API. When the developer clicks Commit in any commit panel:

1. **Detect** — `beforeCheckin()` checks if any files in the commit are tracked by `WorkingSetService`
2. **Separate** — AI-tracked files are extracted from the commit set
3. **Commit** — `CopilotCommitService` writes the agent's known content to disk, runs `git commit --author="GitHub Copilot (model) <copilot@copilot.example>" -- <paths>` with the `Generated-by` trailer, then restores the original disk content
4. **Notify** — A balloon notification reports the result: "Committed N file(s) as GitHub Copilot"
5. **Continue** — If non-AI files remain, IntelliJ's normal commit proceeds for those files under the developer's name

This write-stage-restore approach works around the `CheckinHandler` API limitation (can only return `COMMIT/CANCEL`, cannot modify author or message) by executing a separate `git commit` for the AI files before IntelliJ's commit runs.

Return value logic

Scenario	Handler returns	Effect
No AI files in commit	COMMIT	Normal IntelliJ commit
AI files committed, non-AI files	COMMIT	AI files committed as Copilot; IntelliJ commits the rest under

Scenario	Handler returns	Effect
remain		developer's name
AI files committed, no remaining files	CANCEL	AI files committed as Copilot; commit dialog closes
Attribution commit fails	COMMIT	Graceful fallback — all files commit under developer's name with warning

Hunk-level diff tracking

HunkDiffEngine uses IntelliJ's `ComparisonManager.compareLines()` API to compute which line ranges were modified by the agent. This enables:

- Logging which specific hunks are agent-authored in each commit
- Future per-line attribution for mixed files (files with both AI and developer edits)

For new files, the entire file is treated as a single agent hunk. For modified files, only the changed line ranges are attributed to the agent.

Worktree isolation for subagents

When a subagent has `forkContext: true` in its `.agent.md` definition, it operates in a git worktree rather than the main workspace. This is primarily a **parallel agent safety** mechanism — multiple subagents can't overwrite each other's changes — but it also adds an approval gate before changes enter the attribution pipeline.

The subagent works in `.copilot-chat/worktrees/<agentId>/` on branch `copilot-worktree-<agentId>`. Its tool calls are routed to the worktree via a per-conversation workspace override in `ToolRouter`. When the subagent completes, a diff is generated against the main workspace and a review panel appears inline in the Agent chat. The user can:

- **Apply** — files are copied to the main workspace and registered with `WorkingSetService`, entering the normal attribution pipeline (Copilot Changes changelist → auto-attributed commit)
- **Discard** — the worktree is deleted with no effect on the main workspace

If no files changed, the worktree is cleaned up automatically. Stale worktrees from crashes are pruned on startup.

Change source summary

Change source	Approval gate	Commit author	Trailer
Developer edits	None (normal flow)	Developer	None
Agent tools (direct)	Changes tab review	GitHub Copilot (model)	Generated-by

Change source	Approval gate	Commit author	Trailer
Agent tools (worktree)	Worktree review → Changes tab	GitHub Copilot (model)	Generated-by

Both agent paths produce identical commits. The worktree path adds a pre-review step before changes reach the Working Set. Attribution is enforced automatically on all commit paths — the developer does not need to use a specific button or workflow.

Querying AI-generated commits

```
# All Copilot commits
git log --author="GitHub Copilot"

# All commits with the Generated-by trailer
git log --grep="Generated-by: github-copilot"

# Count of Copilot vs developer commits
echo "Copilot: $(git log --author='GitHub Copilot' --oneline | wc -l)"
echo "Developer: $(git log --author='GitHub Copilot' --invert-grep --oneline | wc -l)"
```

On GitHub, the `--author` field shows in the commit detail view and PR commit list. The `Generated-by` trailer renders as structured metadata below the commit message. PR reviewers can see at a glance which commits were AI-authored.

Comparison with alternatives

Approach	Attribution in git log	GitHub UI	Requires account	Parseable by CI
<code>--author + trailer</code> (ours)	Author field + trailer	Trailer in commit view	No	Yes
Co-authored-by trailer	Trailer only	Co-author avatar	No	Yes
GitHub machine user	Real author with avatar	Full profile link	Yes (1 seat)	Yes
GitHub App	app[bot] author	Bot badge	Yes (app setup)	Yes
Commit message prefix [copilot]	Message only	Visible but unstructured	No	Fragile regex

Industry precedent

- Claude Code appends `Co-Authored-By: Claude <noreply@anthropic.com>` to commits

- **GitHub Copilot coding agent** commits as `copilot[bot]` (server-side only)
- **Cursor** relies on manual attribution via `.cursorrules` configuration
- **SSW Rules** recommends co-author trailers for all AI-assisted commits

The `--author` + trailer approach provides stronger attribution than co-author trailers (Copilot is the *author*, not a co-author) while remaining fully local with no account dependencies.